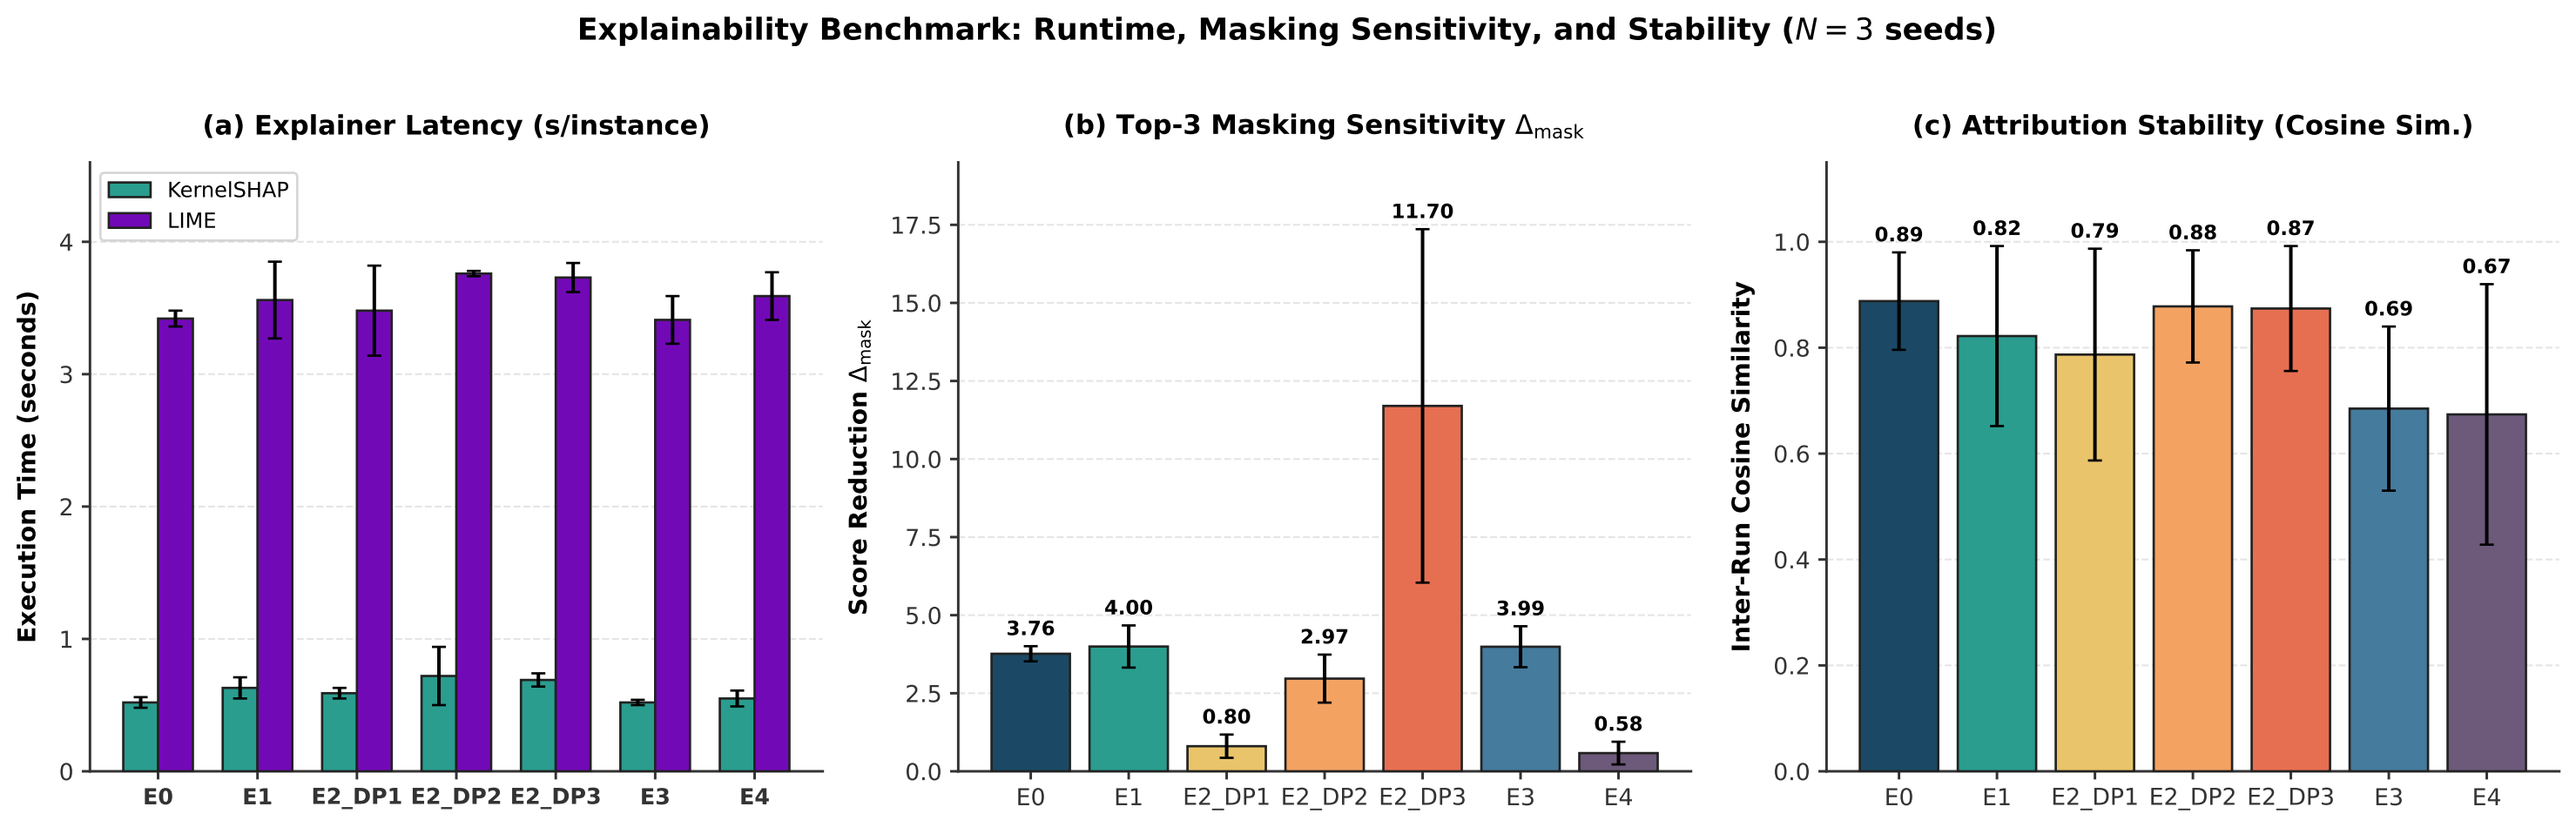
Explainability Benchmark: Runtime, Masking Sensitivity, and Stability ($N = 3$ seeds)
(a) Explainer Latency (s/instance)
KernelSHAP
LIME
Execution Time (seconds)
0
1
2
3
4
E0
E1
E2_DP1
E2_DP2
E2_DP3
E3
E4
(b) Top-3 Masking Sensitivity Δ_{mask}
Score Reduction Δ_{mask}
0.0
2.5
5.0
7.5
10.0
12.5
15.0
17.5
3.76
4.00
0.80
2.97
11.70
3.99
0.58
E0
E1
E2_DP1
E2_DP2
E2_DP3
E3
E4
(c) Attribution Stability (Cosine Sim.)
Inter-Run Cosine Similarity
0.0
0.2
0.4
0.6
0.8
1.0
0.89
0.82
0.79
0.88
0.87
0.69
0.67
E0
E1
E2_DP1
E2_DP2
E2_DP3
E3
E4